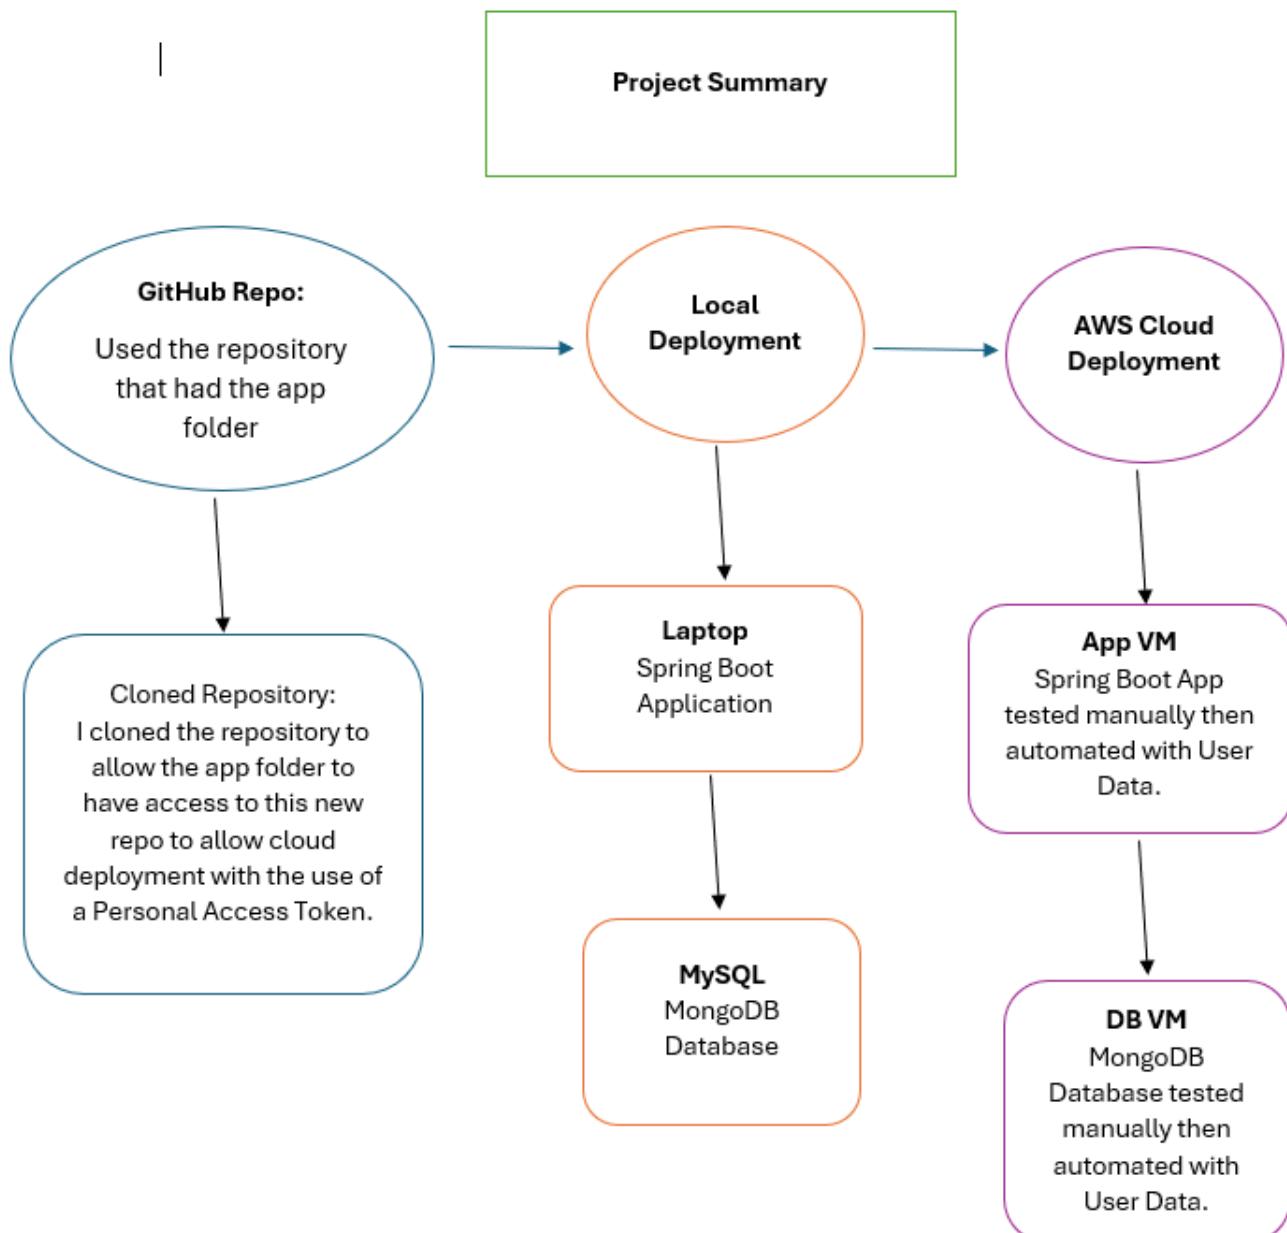


- Java Springboot app Notes
 - Project Architecture
- Step 1: Create a Git repo
- Step 2: Deploy the app locally
- Blockers/Errors:
- Step 3: Deploying app using AWS VM
- DB VM:
- APP VM:
- Blockers/Errors:
- Step 4: Deploy app and database on AWS using Docker
- Step 5: Deploy app and database on AWS using Kubernetes (Minikube)
- What I learnt from this project

Java Springboot app Notes

Project Architecture



This shows the overall project architecture and the steps I have take to get the app deployed using the cloud with applications and platforms such as Kubernetes and Docker.

It begins with a GitHub repository containing the application folder. That repository is then cloned into a new repo through a PAT (Personal Access Token) to ensure that the application folder is accessible for later cloud deployment steps. From there, the workflow proceeds to Local Deployment, where the Spring Boot application runs on a laptop and connects to a local MySQL/MongoDB database for initial testing. Once the application is verified locally, the process moves to AWS Cloud Deployment, where two virtual machines (VMs) are used: one for running the Spring Boot application (App VM) and another for the database (DB VM). Both VMs are tested manually at first and then automated using User Data scripts, enabling consistent, repeatable deployment in the cloud.

Step 1: Create a Git repo

- Create a Git repo and clone it into Git bash.
- Follow the steps from the README and make sure the repo is private.
- Commands to follow:
 - `git init`
 - `git add .`
 - `git commit -m "Initial commit of file"`
 - `git remote add origin "https://github.com/<USERNAME>/repo-name"`
 - `git branch -M main`
 - `git push -u origin main`

Step 2: Deploy the app locally

- In order to get the app working locally, we need to install 3 things:
 - Java 17
 - Maven
 - MySQL
- After installing the right dependencies and verifying by checking the version in the terminal, login to MySQL:
 - Connect manually to root user: `\connect root@localhost`
 - Make sure you're in MySQL Shell by typing `\sql`
 - `mysql -u root -p`
- Set password to `Burntwood1!`
- We would need to create the database first as we haven't made it before nor has it been made for us.
 - CREATE DATABASE library;
 - USE library;
- Run this code in the terminal after making sure that you're in the app's repo where the file `library.sql` is. This is what will seed the data in the database.
- In the terminal, run: `mysql -u root -p library < library.sql`
- Verify the tables and data by entering this in the MySQL Command line:
 - `'SHOW TABLES;`
 - `SELECT * FROM authors;`
- Once you can see the data, we need to enter the environmental variables.

- The app connects to MySQL using `DB_HOST`, `DB_USER` and `DB_PASS`. I set this in my local machine using the "Windows" button to "edit the system environment variables".
- `DB_HOST`: `jdbc:mysql://localhost:3306/library`
- `DB_USER`: `root`
- `DB_PASS`: `Burntwood1!`
- After setting the variables, I had to restart the terminal again to make sure that the variables apply.
- After entering the terminal, I made sure to get my file path to the app's directory. Once I'm in the repo I did `cd LibraryProject2`.
- Then I used Maven to start the app: `mvn spring-boot:run`
- After the process runs, look out for these lines: `Tomcat started on port(s): 5000 (http)`
`Started SpringapiApplication in 3.234 seconds`
- To test the app, in the browser go to: `http://localhost:5000/web/authors`. This will show the list of the 4 authors from the database.
- To see all authors: `http://localhost:5000/authors`.
- Single author (ID 3): `http://localhost:5000/author/3`
- To stop the app: `mvn spring-boot:stop`

Blockers/Errors:

- I had the longest time trying to fix the app using Java V17 even after installing the correct one.
- The app was continuously using Java V1 and I couldn't understand why.
- I just opened a new terminal and checked the version that the app was to use after doing `java -version` and when it said 17, I understood that it now works.
- Also, I created a new environmental variable named "JAVA_HOME" and set it to JDK's path.
- Then under "System variables", I found "Path" and created a new variable: `%JAVA_HOME%\bin`.

Step 3: Deploying app using AWS VM

- I created 2 VM. One for DB and one for the app.

DB VM:

- Ubuntu 22.04 EC2 instance with t3.micro image (free tier).
- Key pair: tech511-afsheen-aws
- Security Groups:
 - Allow SSH (port 22)
 - Add Security Group rule:
 - Type: MYSQL/Aurora
 - Port: 3306
 - Source: Custom
 - Enter: 0.0.0.0/0 (for testing)
 - Add User Data
- After launching the instance, I SSH'd into a new terminal and typed in `sudo systemctl status mysql`.
- This showed "active" which told me that everything was working well.

APP VM:

- Ubuntu 22.04 EC2 instance with t3.micro image (free tier)
- Key pair: tech511-afsheen-aws
- Security Groups:
 - Allow SSH (port 22)
 - Custom TCP (port 5000) (for app web access)
- Add User Data but make sure that you type in the Private IPv4 address from the DB VM.
- From the User Data, enter your github username and repo name.
- To verify everything works, got to the app VM and be in the java app repo where `pom.xml` is. Click on the public IPv4 address from the app VM: `http://<APP_PUBLIC_IP>:5000/web/authors`
- Make sure you have cloned the repo in the DB VM. I had not done this and had to do a `git clone` using a new SSH key token since the repo is private and I couldn't remember the username and password it was asking for.
- In the App VM, to build the app, I did this: `./mvnw clean package -DskipTests`

Blockers/Errors:

- I had problem with SQL rejecting permissions for the app to access the database even though it could. It only gave access for 'root' to access the db.
- I had to change the environmental variables in the bash scripts/user data that changed the permissions to allow the app to access the DB.

- Also, I kept on forgetting to get to the correct file path that wasn't getting the app deployed.

```
ubuntu@ip-172-31-62-112:~$ mvn spring-boot:run
[INFO] Scanning for projects...
Downloading from central: https://repo.maven.apache.org/maven2/org/apache/maven/plugins/maven-metadata.xml
Downloading from central: https://repo.maven.apache.org/maven2/org/codehaus/mojo/maven-metadata.xml
Downloaded from central: https://repo.maven.apache.org/maven2/org/codehaus/mojo/maven-metadata.xml (21 kB at 29 kB/s)
Downloaded from central: https://repo.maven.apache.org/maven2/org/apache/maven/plugins/maven-metadata.xml (14 kB at 20 kB/s)
[INFO] -----
[INFO] BUILD FAILURE
[INFO] -----
[INFO] Total time: 1.821 s
[INFO] Finished at: 2025-10-21T13:48:21Z
[INFO] -----
[ERROR] No plugin found for prefix 'spring-boot' in the current project and in the plugin groups [org.apache.maven.plugins, org.codehaus.mojo] available from the repositories [local (/home/ubuntu/.m2/repository), central (https://repo.maven.apache.org/maven2)] -> [Help 1]
[ERROR]
[ERROR] To see the full stack trace of the errors, re-run Maven with the -e switch.
[ERROR] Re-run Maven using the -X switch to enable full debug logging.
[ERROR]
[ERROR] For more information about the errors and possible solutions, please read the following articles:
[ERROR] [Help 1] http://cwiki.apache.org/confluence/display/MAVEN/NoPluginFoundForPrefixException
ubuntu@ip-172-31-62-112:~$ ^C
ubuntu@ip-172-31-62-112:~$ ls
library-java17-mysql-app
ubuntu@ip-172-31-62-112:~$ cd library-java17-mysql-app
ubuntu@ip-172-31-62-112:~/library-java17-mysql-app$ ls
LibraryProject2 README.md library.sql readme-images
ubuntu@ip-172-31-62-112:~/library-java17-mysql-app$ cd LibraryProject2
ubuntu@ip-172-31-62-112:~/library-java17-mysql-app/LibraryProject2$ mvn spring-boot:run
```

- One error I got was that MySQL was not recognising the environment variables I set and constantly asked for variables to start with JDBC. To fix this I had to set the environmental variables again:

- `export SPRING_DATASOURCE_URL=jdbc:mysql://172.31.25.143:3306/library`
- `export SPRING_DATASOURCE_USERNAME=root`
- `export SPRING_DATASOURCE_PASSWORD=Burntwood1!`

```
2025-10-21T14:49:36.825Z ERROR 1142 --- [ restartedMain] j.LocalContainerEntityManagerFactoryBean : Failed to initialize JPA EntityManagerFactory: Unable to create requested service [org.hibernate.engine.jdbc.env.spi.JdbcEnvironment] due to: Unable to determine Dialect without JDBC metadata (please set 'javax.persistence.jdbc.url', 'hibernate.connection.url', or 'hibernate.dialect')
2025-10-21T14:49:36.826Z WARN 1142 --- [ restartedMain] ConfigServletWebServerApplicationContext : Exception encountered during context initialization - cancelling refresh attempt: org.springframework.beans.factory.BeanCreationException: Error creating bean with name 'entityManagerFactory' defined in class path resource [org/springframework/boot/autoconfigure/orm/jpa/HibernateJpaConfiguration.class]: Unable to create requested service [org.hibernate.engine.jdbc.env.spi.JdbcEnvironment] due to: Unable to determine Dialect without JDBC metadata (please set 'javax.persistence.jdbc.url', 'hibernate.connection.url', or 'hibernate.dialect')
2025-10-21T14:49:36.830Z INFO 1142 --- [ restartedMain] o.apache.catalina.core.StandardService : Stopping service [Tomcat]
2025-10-21T14:49:36.845Z INFO 1142 --- [ restartedMain] .s.b.a.l.ConditionEvaluationReportLogger :
Error starting ApplicationContext. To display the condition evaluation report re-run your application with 'debug' enabled.
2025-10-21T14:49:36.873Z ERROR 1142 --- [ restartedMain] o.s.boot.SpringApplication : Application run failed
```

Step 4: Deploy app and database on AWS using Docker

- To start of this task, I went on to create an EC2 instance on AWS using an **22.04 LTS AMI**, a **t3.small** instance type and inbound ports 22, 80, 3000, and 27017.
- I then added user data:

```
#!/bin/bash
apt-get update -y
apt-get install -y git docker.io docker-compose -y
systemctl enable docker
systemctl start docker

# Deploy app automatically
cd /home/ubuntu
git clone https://github.com/Afsheen28/tech511-sparta-app.git
cd tech511-sparta-app
docker-compose up -d
```

- Once the instance had launched, I SSH'd into the VM.
- I then cloned my **app** folder from my **tech511-sparta-app** to this **java-spring-boot-app-notes** repository. This is because the app folder was not in this repo.

```
git clone https://github.com/Afsheen28/java-spring-boot-app-notes.git
Cloning into 'java-spring-boot-app-notes'...
remote: Enumerating objects: 3, done.
remote: Counting objects: 100% (3/3), done.
remote: Total 3 (delta 0), reused 3 (delta 0), pack-reused 0 (from 0)
Receiving objects: 100% (3/3), done.
ubuntu@ip-172-31-32-24:~$ ls
java-spring-boot-app-notes  tech511-sparta-app
ubuntu@ip-172-31-32-24:~$ cp -r ~/tech511-sparta-app/app/* ~/java-spring-boot-a
pp-notes/
ubuntu@ip-172-31-32-24:~$ cd java-spring-boot-app-notes/
ubuntu@ip-172-31-32-24:~/java-spring-boot-app-notes$ ls -la
total 44
drwxrwxr-x 8 ubuntu docker 4096 Nov 12 12:34 .
drwxr-x--- 6 ubuntu ubuntu 4096 Nov 12 12:33 ..
drwxrwxr-x 8 ubuntu docker 4096 Nov 12 12:33 .git
-rw-rw-r-- 1 ubuntu docker 1681 Nov 12 12:34 README.md
-rw-r--r-- 1 ubuntu docker 1415 Nov 12 12:34 app.js
drwxr-xr-x 2 ubuntu docker 4096 Nov 12 12:34 models
-rw-r--r-- 1 ubuntu docker 655 Nov 12 12:34 package.json
drwxr-xr-x 3 ubuntu docker 4096 Nov 12 12:34 public
drwxr-xr-x 2 ubuntu docker 4096 Nov 12 12:34 seeds
drwxr-xr-x 2 ubuntu docker 4096 Nov 12 12:34 test
drwxr-xr-x 5 ubuntu docker 4096 Nov 12 12:34 views
```

- Once the app folder was in the repo, I created a `docker-compose.yml` and `provision.sh` file.

```
version: '3.8'

services:
  app:
    image: node:20
    container_name: node_app
    working_dir: /usr/src/app
    volumes:
      - ./usr/src/app
    ports:
      - "3000:3000"
    environment:
      - DB_HOST=mongodb://mongo:27017/posts
    depends_on:
      - mongo
    command: bash -c "npm install && node app.js"

  mongo:
    image: mongo:6
    container_name: mongo_db
    restart: always
    ports:
      - "27017:27017"
    volumes:
      - mongo_data:/data/db

volumes:
  mongo_data:
```

```

#!/bin/bash
# Update and install Docker
sudo apt-get update -y
sudo apt-get install -y docker.io docker-compose

# Enable and start Docker
sudo systemctl enable docker
sudo systemctl start docker

# Clone the repo or copy files
cd /home/ubuntu
git clone https://github.com/Afsheen28/tech511-sparta-app.git
cd tech511-sparta-app

# Bring up Docker containers
sudo docker-compose up -d

# Optional: Seed the database
sudo docker exec -it node_app bash -c "node seeds/seed.js"

echo "Deployment complete! App is running on port 3000"

```

- I then updated the system and installed the required packages (like docker):

```

ubuntu@ip-172-31-32-24:~$ sudo apt update -y
Hit:1 http://eu-west-1.ec2.archive.ubuntu.com/ubuntu jammy InRelease
Hit:2 http://eu-west-1.ec2.archive.ubuntu.com/ubuntu jammy-updates InRelease
Hit:3 http://eu-west-1.ec2.archive.ubuntu.com/ubuntu jammy-backports InRelease
Hit:4 http://security.ubuntu.com/ubuntu jammy-security InRelease
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
18 packages can be upgraded. Run 'apt list --upgradable' to see them.
ubuntu@ip-172-31-32-24:~$ sudo apt install -y docker.io docker-compose git
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
docker-compose is already the newest version (1.29.2-1).
git is already the newest version (1:2.34.1-1ubuntu1.15).
docker.io is already the newest version (28.2.2-0ubuntu1~22.04.1).
0 upgraded, 0 newly installed, 0 to remove and 18 not upgraded.
ubuntu@ip-172-31-32-24:~$ sudo systemctl start docker
ubuntu@ip-172-31-32-24:~$ sudo usermod -aG docker $USER

```

- I then gave user permissions to run Docker without `sudo`. By default only root can access Docker.

Adding the `ubuntu` user to the docker group allows Docker commands without `sudo`.

```

ubuntu@ip-172-31-32-24:~$ sudo usermod -aG docker $USER
ubuntu@ip-172-31-32-24:~$ newgrp docker
ubuntu@ip-172-31-32-24:~$ docker ps

```

- | COUNTAINER ID | IMAGE | COMMAND | CREATED | STATUS | PORTS | NAMES |
|---------------|-------|---------|---------|--------|-------|-------|
|---------------|-------|---------|---------|--------|-------|-------|

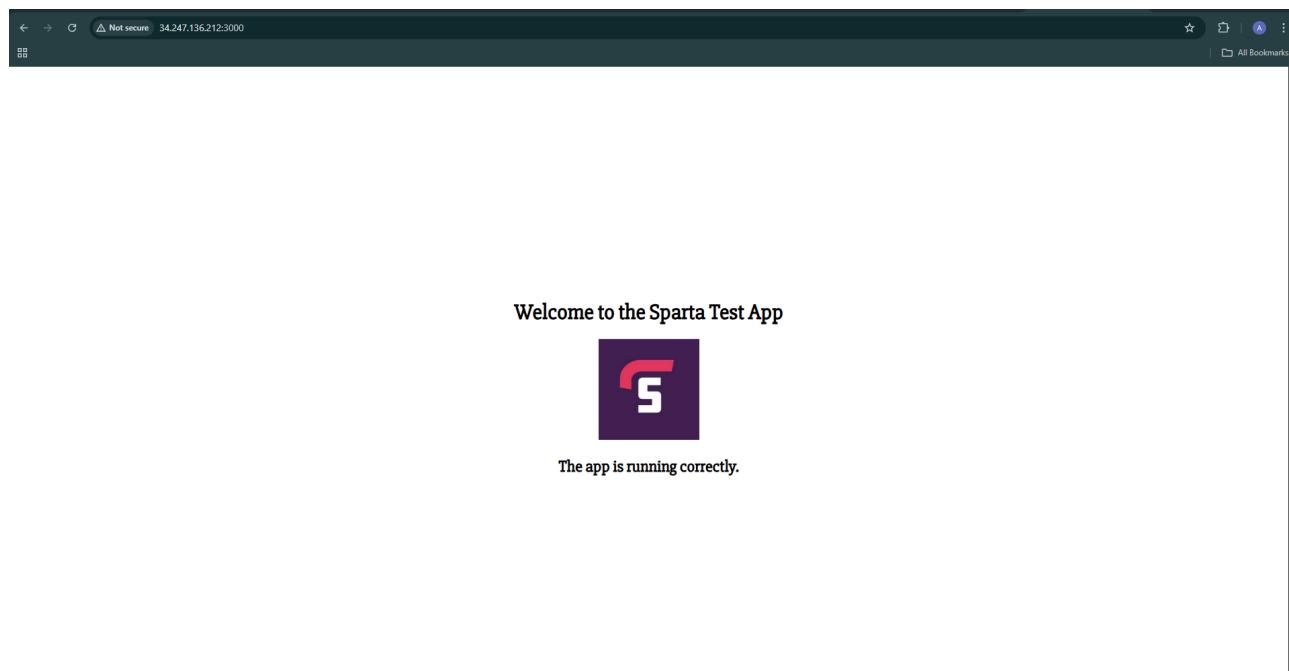
- I then verified docker access by running `docker ps`. This confirmed that Docker was working and accessible to the user. An empty table was shown meaning that no containers were running yet which

was true.

- I then deployed the containers by running `docker-compose up -d`. This launched both the Node.js and MongoDB containers in detached mode. Compose automatically created an internal network linking the app to the database.
- After that, I ran `docker ps` to verify both services were running and mapped to their ports.

```
ubuntu@ip-172-31-32-24:~/java-spring-boot-app-notes$ docker-compose up -d
Creating network "java-spring-boot-app-notes_default" with the default driver
Creating volume "java-spring-boot-app-notes_mongo_data" with default driver
Pulling mongo (mongo:6)...
6: Pulling from library/mongo
af6eca94c810: Pull complete
1ad79419bb76: Pull complete
afdef96060ff: Pull complete
dc9abaf2c8ff: Pull complete
a2be9b1fc57a: Pull complete
9d81a4e85c6a: Pull complete
9d6e74542339: Pull complete
cde22982277a: Pull complete
Digest: sha256:0a83b2f824b29c94c5284a49b018db2e3166d9879d1edda5cce545d899ef2bbb
Status: Downloaded newer image for mongo:6
Pulling app (node:20)...
20: Pulling from library/node
5d93aea69798: Pull complete
bb445e472b1b: Pull complete
2123190679e8: Pull complete
32885a2b0a58: Pull complete
fb9baa9d1d1d: Pull complete
e6819020f277: Pull complete
de002888bed8: Pull complete
b82a1e14a32d: Pull complete
Digest: sha256:47dacd49500971c0fbe602323b2d04f6df40a933b123889636fc1f76bf69f58a
Status: Downloaded newer image for node:20
Creating mongo_db ... done
Creating node_app ... done
ubuntu@ip-172-31-32-24:~/java-spring-boot-app-notes$ docker ps
CONTAINER ID   IMAGE       COMMAND                  CREATED             STATUS              NAMES
           PORTS
31e0ccb89c48   node:20    "docker-entrypoint.s..."  9 seconds ago      Up 8 seconds        node_app
               0.0.0.0:3000->3000/tcp, [::]:3000->3000/tcp
bfbaf18cf6a4    mongo:6    "docker-entrypoint.s..."  10 seconds ago     Up 8 seconds        mongo_db
               0.0.0.0:27017->27017/tcp, [::]:27017->27017/tcp
```

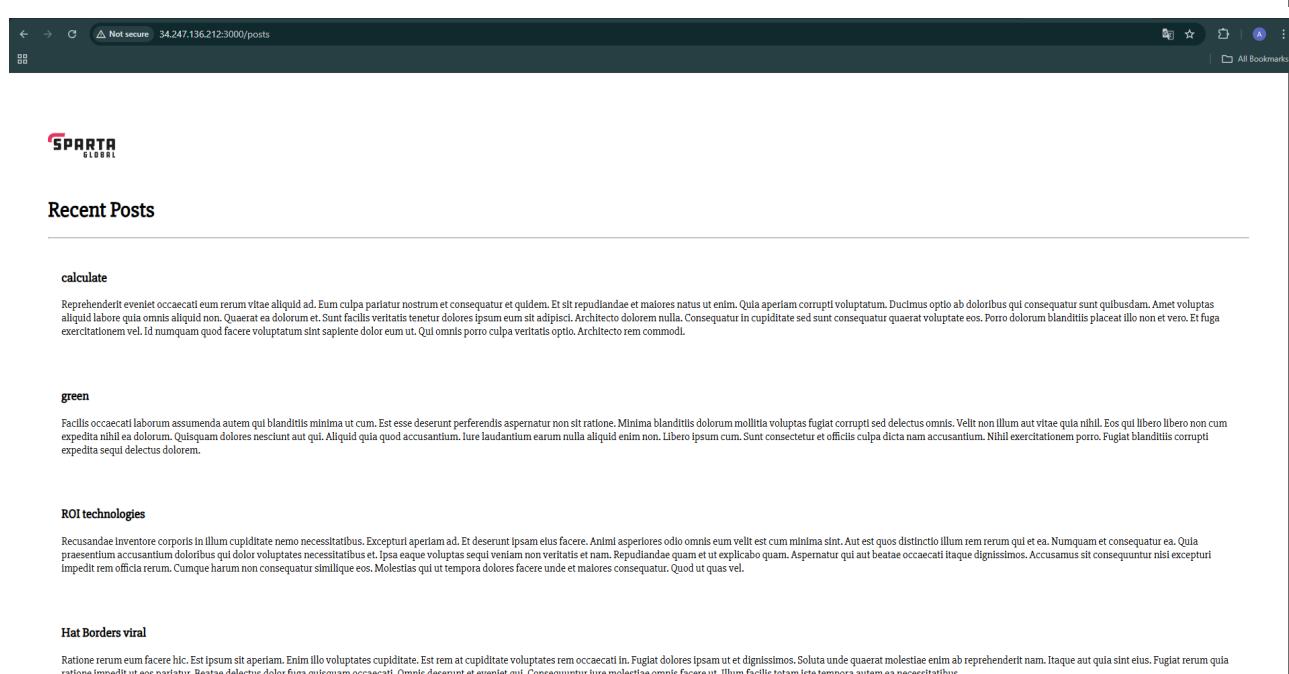
- I then tested the application by going on `http://<EC2_PUBLIC_ID>:3000` and `http://<EC2_PUBLIC_ID>:3000/posts`



Welcome to the Sparta Test App



The app is running correctly.



Recent Posts

calculate
Reprehenderit eveniet occaecati eum rerum vitae aliquid ad. Eum culpa paratur nostrum et consequatur et quidem. Et sit repudiandae et maiores natus ut enim. Quia aperiam corrupti voluptatum. Ducimus optio ab doloribus qui consequatur sunt quibusdam. Amet voluptas aliquid labore quia omnis aliquid non. Quaerat ea dolorum et. Sunt facilis veritatis tenetur dolores ipsum eum sit adipisci. Architecto dolorem nulla. Consequatur in cupiditate sed sunt consequatur quaerat voluptate eos. Porro dolorum blanditiis placeat illo non et vero. Et fuga exercitationem vel. Id numquam quod facere voluptatum sint sapiente dolor eum ut. Qui omnis porro culpa veritatis optio. Architecto rem commodi.

green
Facilis occaecati laborum assumenda autem qui blanditiis minima ut cum. Est esse deserunt preferendis aspernatur non sit ratione. Minima blanditiis dolorum mollitia voluptas fugiat corrupti sed delectus omnis. Velt non illum aut vitae quia nihil. Eos qui libero libero non cum expedita nihil ea dolorum. Quisquam dolores nesciunt aut qui. Aliquid quia quod accusantium. Iure laudantium earum nulla aliquid enim non. Libero ipsum cum. Sunt consectetur et officiis culpa dicta nam accusantium. Nihil exercitationem porro. Fugiat blanditiis corrupti expedita sequi delectus dolorem.

ROI technologies
Recusandae inventore corporis in illum cupiditate nemo necessitatibus. Excepturi aperiam ad. Et deserunt ipsam eius facere. Animi asperiores odio omnis eum velit est cum minima sint. Aut est quos distinctio illum rem rerum qui et ea. Numquam et consequatur ea. Quia praesentium accusantium doloribus qui dolor voluptates necessitatibus et. Ipsa eaque voluptas sequi veniam non veritatis et nam. Repudiandae quam et ut explicabo quam. Aspernatur qui aut beatae occaecati itaque dignissimos. Accusamus sit consequuntur nisi excepturi impedit rem officia rerum. Cumque harum non consequatur similique eos. Molestias qui ut tempora dolores facere unde et maiores consequatur. Quod ut quas vel.

Hat Borders viral
Ratione rerum eum facere hic. Est ipsum sit speriam. Enim illo voluptates cupiditate. Est rem at cupiditate voluptates rem occaecati in. Fugiat dolores ipsam ut et dignissimos. Soluta unde quaerat molestiae enim ab reprehenderit nam. Itaque aut quia sint eius. Fugiat rerum quia ratione impedit ut eos paratur. Beatae delectus dolor fuga quisquam occaecati. Omnis deserunt et eveniet qui. Consequuntur iure molestiae omnis facere ut. Illum facilis totam iste tempora autem ea necessitatibus.

Step 5: Deploy app and database on AWS using Kubernetes (Minikube)

- Firstly, I created an EC2 instance with an instance type of t3.medium due to the RAM size I needed. I also added the following inbound rules in the security group.

The screenshot shows a web-based interface for managing network security rules. At the top, there are various icons and a header bar with the text "sparta-devops (1359-2847-6890)" and "tech511-afsheen". Below the header, a search bar contains the identifier "31fde6da47c6". The main area is divided into two sections: "Inbound rules" and "Outbound rules".

Inbound rules:

Name	Security group rule ID	Port range	Protocol	Source
-	sgr-0642abb72e9a8275d	3000	TCP	0.0.0.0/0
-	2 IDs	22	TCP	0.0.0.0/0
-	2 IDs	80	TCP	0.0.0.0/0
-	sgr-090d642a89701f5d4	30000 - 32767	TCP	0.0.0.0/0
-	sgr-08b81beee8a7e5db9	443	TCP	0.0.0.0/0
-	sgr-002b31cacc8f7a62a	30002	TCP	0.0.0.0/0

Outbound rules:

- I then added user data:

```
$ user-data-kubernetes.sh
1  #!/bin/bash
2  set -e
3
4  # --- Update System ---
5  sudo apt update -y && sudo apt upgrade -y
6
7  # --- Install Dependencies ---
8  sudo apt install -y apt-transport-https ca-certificates curl software-properties-common conntrack
9
10 # --- Install Docker ---
11 sudo apt install -y docker.io
12 sudo systemctl enable docker
13 sudo systemctl start docker
14
15 # --- Install Minikube ---
16 curl -Lo minikube https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
17 sudo install minikube /usr/local/bin/minikube
18
19 # --- Install Kubectl ---
20 curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
21 sudo install kubectl /usr/local/bin/kubectl
22
23 # --- Create Working Directory ---
24 mkdir -p /home/ubuntu/k8s-deploy
25 cd /home/ubuntu/k8s-deploy
26
27 # --- Start Minikube ---
28 sudo minikube start --driver=docker --force
29
30 # --- Apply YAML Files ---
31 sudo kubectl apply -f mongo-pv-pvc.yml
32 sudo kubectl apply -f mongo-deployment.yml
33 sudo kubectl apply -f nodejs-deployment.yml
34
35 # --- Wait for Pods ---
36 sudo kubectl wait --for=condition=Ready pods --all --timeout=300s
37
```

```
# --- Wait for Pods ---
sudo kubectl wait --for=condition=Ready pods --all --timeout=300s

# --- Display Deployment Info ---
sudo kubectl get pods -o wide
sudo kubectl get svc -o wide
```

- Once I launched the instance, and SSH'd into a new Gitbash terminal, I created a repository named `k8s-deploy` and created 3 YAML files.

```
ubuntu@ip-172-31-54-133:~/k8s-deploy$ cat mongo-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mongo-deploy
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mongo
  template:
    metadata:
      labels:
        app: mongo
    spec:
      containers:
        - name: mongo
          image: mongo:6
          ports:
            - containerPort: 27017
          volumeMounts:
            - name: mongo-storage
              mountPath: /data/db
        volumes:
          - name: mongo-storage
            persistentVolumeClaim:
              claimName: mongo-pvc
---
apiVersion: v1
kind: Service
metadata:
  name: mongo-service
spec:
  selector:
    app: mongo
  ports:
    - port: 27017
      targetPort: 27017
  clusterIP: None
```

```
ubuntu@ip-172-31-54-133:~/k8s-deploy$ cat mongo-pv-pvc.yaml
apiVersion: v1
kind: PersistentVolume
metadata:
  name: mongo-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: /mnt/data/mongo
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mongo-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
```

```
ubuntu@ip-172-31-54-133:~/k8s-deploy$ cat nodejs-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nodejs-deploy
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nodejs
  template:
    metadata:
      labels:
        app: nodejs
    spec:
      containers:
        - name: nodejs
          image: siddiquia280504/tech511-sparta-app:v1
          env:
            - name: DB_HOST
              value: mongodb://mongo-service:27017/posts
          ports:
            - containerPort: 3000
---
apiVersion: v1
kind: Service
metadata:
  name: nodejs-service
spec:
  selector:
    app: nodejs
  type: NodePort
  ports:
    - port: 3000
      targetPort: 3000
      nodePort: 30002
```

- I then went to start minikube by running `minikube start --driver=docker --memory=2500mb`

```
ubuntu@ip-172-31-54-133:~$ minikube start --driver=docker --memory=2500mb
* minikube v1.37.0 on Ubuntu 22.04
* Using the docker driver based on existing profile
! You cannot change the memory size for an existing minikube cluster. Please first delete the cluster.

X Requested memory allocation (1800MB) is less than the recommended minimum 1900 MB. Deployments may fail.

* Starting "minikube" primary control-plane node in "minikube" cluster
* Pulling base image v0.0.48 ...
* docker "minikube" container is missing, will recreate.
* Creating docker container (CPUs=2, Memory=1800MB) ...
* Preparing Kubernetes v1.34.0 on Docker 28.4.0 ...
* Configuring bridge CNI (Container Networking Interface) ...
* Verifying Kubernetes components...
- Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: storage-provisioner, default-storageclass

! /usr/local/bin/kubectl is version 1.31.0, which may have incompatibilities with Kubernetes 1.34.0.
- Want kubectl v1.34.0? Try 'minikube kubectl -- get pods -A'
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
ubuntu@ip-172-31-54-133:~$ minikube status
minikube
type: Control Plane
host: Running
kubelet: Running
apiserver: Running
kubeconfig: Configured

ubuntu@ip-172-31-54-133:~$ kubectl get nodes
NAME      STATUS    ROLES      AGE      VERSION
minikube  Ready     control-plane   33s    v1.34.0
```

- Minikube had initially warned that the memory allocation (1800MB) was below the recommended limit (1900MB). I then recreated the Minikube cluster with the sufficient memory (2.5GB) and confirmed the status as shown in the picture above.
- I then went ahead and applied the YAML files by running `kubectl apply -f fileName.yml`.

```
ubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl apply -f mongo-pv-pvc.yaml
persistentvolume/mongo-pv created
persistentvolumeclaim/mongo-pvc created
ubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl apply -f mongo-deployment.yaml
deployment.apps/mongo-deploy created
Warning: spec.SessionAffinity is ignored for headless services
service/mongo-service created
ubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl apply -f nodejs-deployment.yaml
deployment.apps/nodejs-deploy created
• service/nodejs-service created
```

- I then ran `kubectl get pods` and `kubectl get svc` to check everything was running.

```
ubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
mongo-deploy-796f87b8fc-2xcnb   1/1     Running   0          73s
nodejs-deploy-69fb6fc9d5-sfvz7   1/1     Running   0          61s
nodejs-deploy-69fb6fc9d5-wr8mc   1/1     Running   0          61s
ubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl get svc
NAME         TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes   ClusterIP   10.96.0.1      <none>           443/TCP      2m4
8s
mongo-service   ClusterIP   None           <none>           27017/TCP    91s
nodejs-service   NodePort    10.103.150.106  <none>           3000:30002/TCP 79s
```

- Both MongoDB and Node.js pods were healthy and responding.

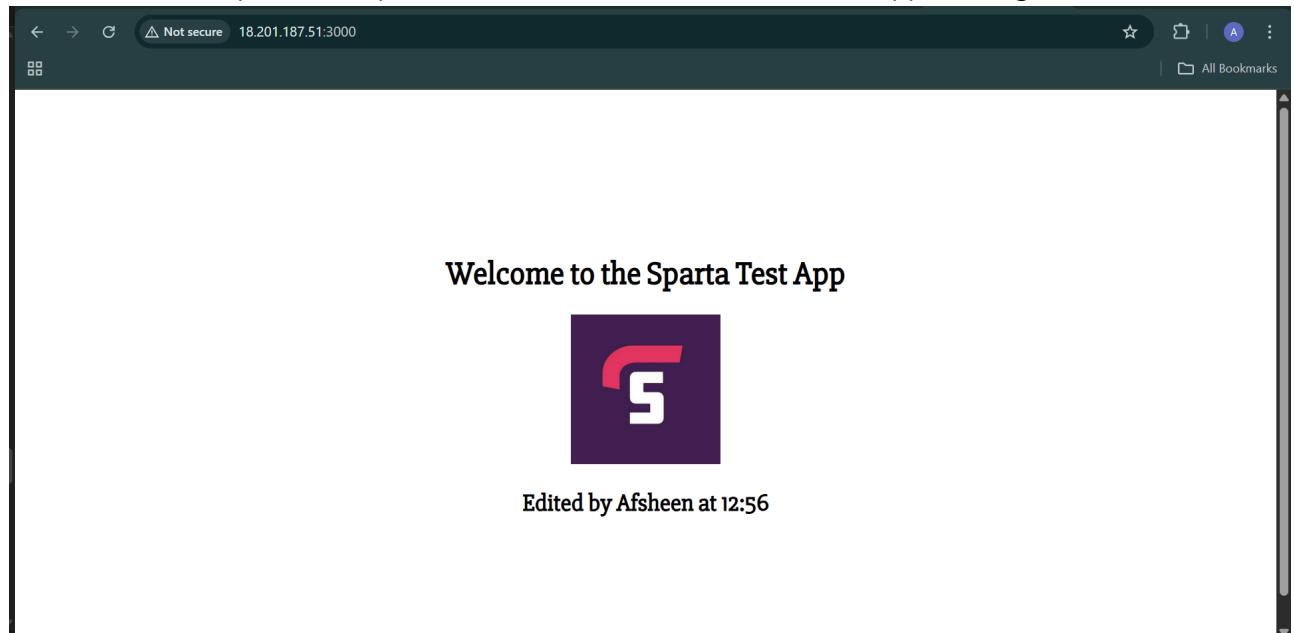
- I then tested the app's internal status

```
ubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl logs -l app=nodejs
Your app is ready and listening on port 3000
Your app is ready and listening on port 3000
```

- Next, I went on to forward the port to 3000 to allow us to see the app running.

```
ubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl port-forward service/nodejs-service 3000:3000 --address=0.0.0.0
Forwarding from 0.0.0.0:3000 -> 3000
Handling connection for 3000
Handling connection for 3000
Handling connection for 3000
Handling connection for 3000
```

- I then went to the public IP + port 3000 of the instance and found the app working.



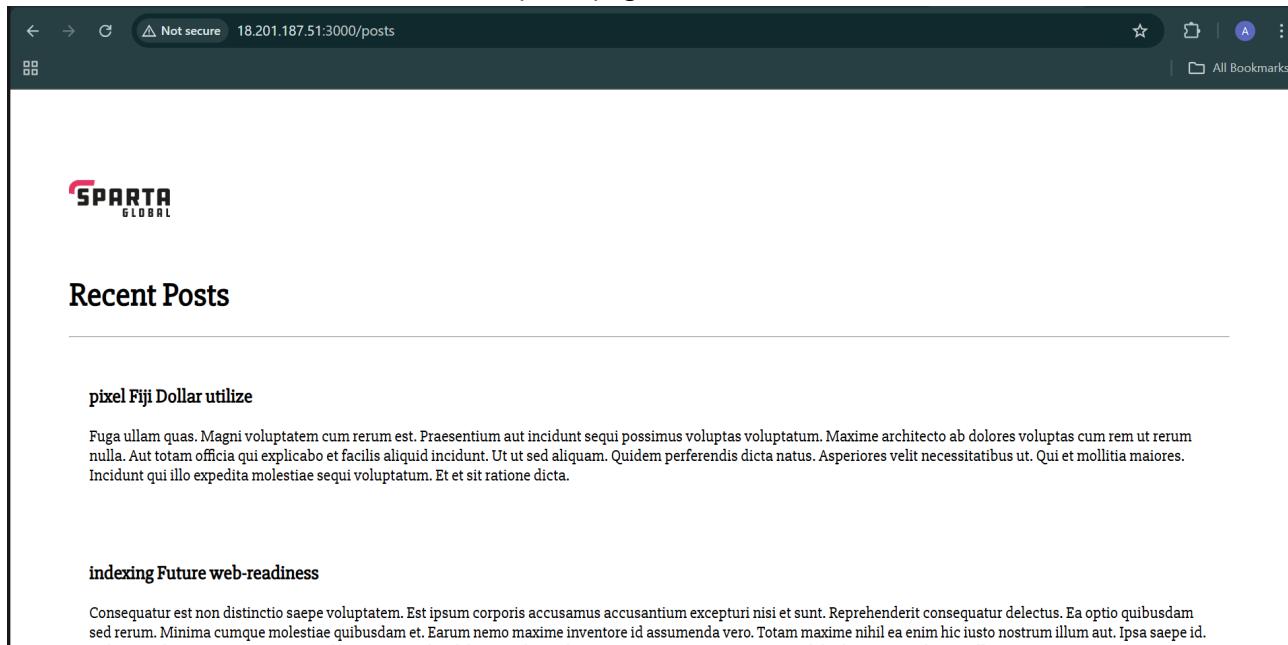
- I then had to seed the database but was not sure what the file path of the `seed.js` was so asked ChatGPT to give me a solution and asked me to run a few commands using one of the Node.js pods where I was then able to seed the database.

```
^Cubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl get pods -l app=nodejs
NAME                      READY   STATUS    RESTARTS   AGE
nodejs-deploy-69fb6fc9d5-sfvz7   1/1     Running   0          11m
nodejs-deploy-69fb6fc9d5-wr8mc   1/1     Running   0          11m
```

```
ubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl exec -it nodejs-deploy-69fb6fc9d5-sfvz7 -- /bin/bash
root@nodejs-deploy-69fb6fc9d5-sfvz7:/usr/src/app# ls /usr/src/app
README.md  models  package-lock.json  public  test
app.js    node_modules  package.json    seeds   views
root@nodejs-deploy-69fb6fc9d5-sfvz7:/usr/src/app# ls /usr/src/app/seeds
seed.js
root@nodejs-deploy-69fb6fc9d5-sfvz7:/usr/src/app# exit
exit
```

```
ubuntu@ip-172-31-54-133:~/k8s-deploy$ kubectl exec -it nodejs-deploy-69fb6fc9d5-sfvz7 -- node /usr/src/app/seeds/seed.js
Connected to database
Database cleared
Database seeded with 100 records
Database connection closed
```

- I then went to the browser and tried the /posts page to check if the database had worked.



What I learnt from this project

During this project, I was able to learn how to deploy an application from end to end using a GitHub repository and SQL seed scripts. Once I started using virtual machines, I was able to use Bash scripts that run through cloud user data, allowing the app and database VMs to configure themselves on startup. I then learned how to containerise my application using Docker, write Dockerfiles, create docker-compose configurations, build and push images to Docker Hub, and deploy them through a single VM provisioning script. I advanced to use Kubernetes to create YAML files to orchestrate the app and database on one machine, with another automated script setting up Kubernetes and applying all configurations. This made me develop my hands-on experience for platforms I had never used before and made me confident on using them as I learnt more about why we use them and what they do along the way.